

Performance Diagnosis for Inefficient Loops

Linhai Song¹ and Shan Lu²


¹FireEye, Inc.

²University of Chicago

What are Performance Problems?

- Definition of Performance Problems (PPs):
 - Implementation mistakes causing inefficiency
- An example


```
char * ssh_xph_generate(node_t * aNode)
{
    int count = 0;
    for (n = aNode; n; n = aNode->prev)
        if(n->name == aNode->name)
            count++;
    return count;
} //Mozilla#477564
```



What are Performance Problems?

- Definition of Performance Problems (PPs):
 - Implementation mistakes causing inefficiency
- An example


```
char * ssh_xph_generate(node_t * aNode)
{
    int count = 0;
    for (n = aNode; n; n = aNode->prev)
        if(n->name == aNode->name)
            count++;
    return count;
} //Mozilla#477564
```



What are Performance Problems?

- Definition of Performance Problems (PPs):
 - Implementation mistakes causing inefficiency
- An example

```
char * ssh_xph_generate(node_t * aNode)
{
    int count = 0;
    for (n = aNode; n; n = aNode->prev)
        if(n->name == aNode->name)
            count++;
    return count;
} //Mozilla#477564
```




Fighting PPs is Important

- PPs widely exist in production run software
 - 5 to 60 Mozilla PPs are fixed each month
- PPs are getting more important
 - Hardware is not getting faster (per-core)
 - Software is getting more complex
 - Energy saving is getting more urgent



Performance Diagnosis

- Identifying the causes of inefficient execution
 - Different from testing & bug detection
 - In-house and on-line diagnosis




Diagnosing PPs is Challenging

- Three criteria for diagnosis tools
 - Coverage: handle a good portion of PPs
 - Accuracy: accurately tell
 - Which code regions are inefficient
 - Why they are inefficient
 - Performance: low run-time overhead





State of the Art

- No existing tools can satisfy the three criteria




Statistical Debugging

[1] Linhai Song, Shan Lu. Statistical Debugging for Real-World Performance Problems. In OOPSLA'2014.




LDoctor

- A tool(kit) targets inefficient loops



Targets of LDoctor



Contributions

- A root-cause taxonomy for inefficient loops
 - Guide the design of automated diagnosis tools
 - Resultless vs. redundant
- LDoctor: a toolkit to diagnose inefficient loops
 - Hybridize static and dynamic analysis
 - Identify inefficiency root causes
 - With high coverage, high accuracy, and low overhead

Outline

- Overview
- Root-cause Taxonomy
- LDoctor Design
- Evaluation of LDoctor
- Conclusion


Outline

- Overview
- Root-cause Taxonomy
- LDoctor Design
- Evaluation of LDoctor
- Conclusion

What Type of Taxonomy We Need?


- Inclusive: covering most/all inefficient loops
- General: Not application specific
- Actionable: helpful to design fix

Root-Cause Taxonomy




Resultless Loops

- Produce no side effects




Resultless 0^*

- Never produce results in any iteration




Resultless $0^*1?$

- Only produce results in the last iteration




Resultless $[0|1]^*$

- May produce results in each iteration




Resultless 1*

- Produce results in almost all iterations




Redundant Loops

- Produce already-available results



Cross-loop Redundancy



Inefficient
Loops

```
char * ssh_xph_generate(node_t * aNode)
{
    int count = 0;
    for (n = aNode; n; n = aNode->prev)
        if(n->name == aNode->name)
            count++;
    return count;
} //Mozilla#477564
```


Buggy Loop

Redundancy


Cross-loop

Redundancy

Cross-iteration Redundancy



Root-Cause Taxonomy





Outline


- Overview
- Root-cause Taxonomy
- LDoctor Design
- Evaluation of LDoctor
- Conclusion

Design Principles


- Taxonomy guided design
- Focused checking
- Static-dynamic hybrid analysis



C.-L. Redundancy Analysis (I)



C.-L. Redundancy Analysis (II)



Outline

- Overview
- Root-cause Taxonomy
- LDoctor Design
- Evaluation of LDoctor
- Conclusion

Implementation & Evaluation

- Implementation using LLVM-3.4.2
- Experimental setting
 - Evaluating 39 bugs from two benchmark suite
 - Applying on look rank list from SD
 - Sampling rate: 1/100 (C.L.) & 1/1000 (C.I.)
- Evaluation metrics



Coverage & Accuracy

BugID	Reported Root Cause	Fix Suggestion	False Positive
Mozilla347306	✓	✓	-
Coverage	✓	✓	-
Mozilla490742	✓	✓	-
Mozilla35294	✓	✓	-
Mozilla477564	✓	✓	-
MySQL27287	✓	✗	01
MySQL15811	✓	✓	-
Apache32546	✓	✓	-
Apache37184	✓	✓	-
Apache29742	✓	✓	-
Apache34464	✓	✓	-
Apache47223	✓	✓	-
...

Coverage & Accuracy

BugID	Reported Root Cause	Fix Suggestion	False Positive
Mozilla347306	✓	✓	-
Mozilla416628	✓	✓	-
Mozilla490742	✓	✓	-
Mozilla35294	✓	✓	-
Mozilla477564	✓	✓	-
MySQL27287	✓	✗	01
MySQL15811	✓	✓	-
Apache32546	✓	✓	-
Apache37184	✓	✓	-
Apache29742	✓	✓	-
Apache34464	✓	✓	-
Apache47223	✓	✓	-
...

Performance

BugID	LDoctor			w/o sampling	
	Resultless	C-L R.	C-I R.	C-L R.	C-I R.
Mozilla347306	1.00	1.18	1.18	356	662
Mozilla416628	1.00	1.04	1.02	73.5	113
MySQL27287	1.05	1.00	-	264	881
MySQL15811	-	1.00	-	414	1088
GCC46401	1.00	1.00	1.00	33.8	60.0
GCC1687	-	/	1.01	/	224
GCC27733	1.00	/	1.00	/	19.4
GCC8805	-	1.00	1.00	3.89	4.36
GCC21430	-	1.04	1.02	165	255
GCC12322	-	1.04	1.00	26.5	25.8

Conclusions

- A root-cause taxonomy for inefficient loops

Our taxonomy are general and specific!


- A series of static-dynamic hybrid analysis

Achieve good coverage, accuracy and performance

Thanks a lot!




Question?







Resultless Analysis




with side effect
without side effect




Loop




0^*



$0^*1?$



$[0|1]^*$



1^*

